


No-Code AI Workflow Pipeline

Flyrank AI Internship — AI Fluency Track | Phase: Build (core) | Task 2: Multi-Step No-Code Workflow

1. Workflow Architecture & Flow Diagram

This pipeline automates technical research, drafting, and quality auditing through a 3-step sequential handoff system.



2. Exact Step Configurations & System Prompts

Step 1: Gather & Synthesize

- **Tool Used:** NotebookLM / Claude
- **Role:** Technical Research Assistant
- **Objective:** Extract core technical facts and remove fluff.

Exact Prompt:

Role: Technical Research Assistant
Task: Analyze the provided raw topic notes/sources below. Extract all core concepts, technical constraints, and primary takeaways.

- Rules:
1. Strip out unnecessary fluff, market hype, or redundant introduction.
 2. Group information logically into key themes.
 3. Output ONLY a structured 4-point outline with clear sub-bullet points.

Input Text:
[Insert Raw Topic / Research Notes Here]

Step 2: First Draft Generation

- **Tool Used:** Claude Project / Custom GPT
- **Role:** Senior Technical Writer
- **Objective:** Convert the structured outline into a comprehensive technical essay.

Exact Prompt:

Role: Senior Technical Writer
Task: Convert the provided 4-point structured outline into a complete, high-quality technical draft.

- Rules:
1. Maintain a professional, clear, and informative technical tone.
 2. Use markdown headings (##, ###) and concise paragraphs.
 3. Do not add speculative information not supported by the outline provided.
 4. Ensure seamless transitions between key sections.

Structured Outline:
[Paste Handoff Output from Step 1]

Step 3: Critique & Revision

- **Tool Used:** ChatGPT (GPT-4o)
- **Role:** Technical Editor & QA Auditor
- **Objective:** Audit the draft against the original outline for accuracy, tone, and missing constraints.

Exact Prompt:

Role: Technical Editor & Quality Assurance Auditor
Task: Critique and revise the draft provided below against the original outline.

Evaluation Criteria:
1. Accuracy: Did the draft miss any core constraints from the outline?
2. Tone & Clarity: Is the phrasing direct, scannable, and free of passive voice/filler?
3. Code/Tech Accuracy: Are framework terms and technical syntax precise?

Instructions:
First, list a brief 2-bullet critique highlighting any weaknesses.
Second, provide the FINAL REVISED VERSION incorporating all corrections.

Original Outline:
[Paste Handoff Output from Step 1]

Draft:
[Paste Handoff Output from Step 2]

3. Live Execution Outputs (Sample Run)

Step 1 Output (NotebookLM Synthesis)

Title: Analysis of Architectural Limits in LLMs and Sensor Fusion in Autonomous Systems

Executive Summary: Synthesizes LLM “Computational Split-Brain Syndrome” (dissociation between explaining rules and executing them due to token embeddings and FFN limits) with autonomous racing robotics (Extended Kalman Filters, LiDAR/RADAR late fusion, and perception delay compensation).

Key Sections:

1. *Architectural Limits:* Contextual contamination, piecewise linear FFNs, instruction-execution disconnect.
2. *Relational Failures:* Decimal comparisons, Alice Problem, tabular reasoning.
3. *Sensor Fusion:* Late fusion EKF, backward-forward delay integration (CTRV model).

Step 2 Output (Claude Draft)

Title: Comprehension Without Competence: Architectural Limits in LLM Reasoning and Lessons from Autonomous Sensor Fusion

Expanded the NotebookLM synthesis into a full 5-part essay covering token embedding failures (e.g., 9.11 vs 9.9), geometric chaos, instruction vs execution pathways, and how safety-critical systems use explicit mathematical frameworks (EKF) instead of emergent pattern matching.

Step 3 Output (ChatGPT Critique Log)

Critique:

- *Accuracy:* Highlighted that Step 2 successfully retained all constraints from Step 1 without introducing hallucinations.
- *Tone & Clarity:* Identified slight sentence verbosity and recommended tighter structural transitions.
- *Pipeline Handoff Note:* Demonstrated self-correction capabilities when a previous step’s text is missing or partially truncated during the handoff.

4. 5-Run Benchmark & Honest Time Accounting

Run #	Input Topic	Manual Time	Pipeline Time	Time Saved
1	LLM Limits vs Autonomous Sensor Fusion	45 mins	8 mins	37 mins
2	Next.js App Router vs Pages Router Architecture	40 mins	7 mins	33 mins
3	Ethical Considerations in Generative AI Systems	35 mins	6 mins	29 mins
4	Firebase Security Rules for Multi-tenant Apps	50 mins	9 mins	41 mins

Time Savings Summary:

- **Total Manual Time (5 runs):** 210 minutes (~3.5 hours)
- **Total Pipeline Time (5 runs):** 37 minutes
- **Initial Setup Overhead:** ~60 minutes (Prompt engineering & workflow testing)
- **Net Time Saved: 113 minutes (~1.9 hours saved, inclusive of setup overhead)**

5. Failure Points & Human Review Checklist

Known Failure Points:

1. **Context/Handoff Loss:** Truncated text during manual copy-pasting between steps causes Step 3 to fail or request missing context.
2. **Niche Syntax Hallucinations:** Step 2 occasionally introduces subtle errors when generating code for brand-new framework updates.
3. **Over-Polishing:** Step 3 can make technical prose sound overly formal, stripping out the author’s personal voice.

Mandatory Human Review Checklist:

- ☐ Verify all numbers, decimal places, and technical parameters against original sources.
- ☐ Check code snippets and framework API syntax manually.
- ☐ Review final tone to ensure alignment with personal/brand publishing guidelines.

Appendix: Full Sample Run Outputs

A.1 Step 1 Output — NotebookLM Synthesis (Full Text)

This briefing document synthesizes critical insights from recent research regarding the fundamental architectural constraints of Large Language Models (LLMs) in symbolic reasoning and the technical implementation of multi-modal sensor fusion for high-speed autonomous racing.

Executive Summary

Contemporary developments in artificial intelligence and autonomous systems reveal a dichotomy between linguistic fluency and execution reliability. Research into LLMs identifies a “Computational Split-Brain Syndrome,” where models can articulate complex procedural rules but fail to execute them due to architectural limitations in symbol representation and feed-forward network (FFN) computation. Specifically, next-token prediction objectives decouple instructional knowledge from execution pathways.

Conversely, in the field of autonomous robotics, reliability is achieved through explicit architectural modularity. The implementation of Extended Kalman Filters (EKF) and multi-modal late fusion (LiDAR and RADAR) allows for high-speed object tracking. A critical innovation in this domain is “delay compensation,” which uses backward-forward integration to reconcile perception latency with real-time motion planning.

I. Architectural Limits of LLMs in Symbolic Reasoning

The “Computational Split-Brain Syndrome” describes a systematic dissociation where LLMs perfectly explain principles they cannot reliably execute. This phenomenon is categorized by three primary architectural constraints.

1. Representation: The Failure of Domain Binding

LLM token embeddings encode context-weighted averages that resist “automatic domain binding.” Because the training objective optimizes for a single embedding across diverse contexts, tokens accumulate conflicting semantic associations.

- *Contextual Contamination:* A token like “9.11” is embedded with associations to historical dates and software versioning. When asked to compare 9.11 to 9.9, the model fails to strip away these non-mathematical associations to treat the token as a pure scalar quantity.
- *Geometric Chaos:* Embedding analysis shows that numerical tokens do not preserve isometric properties (numerical order or distance). For example, “10” may be geometrically closer to “1” than to “9” in certain model latent spaces.

2. Computation: Feed-Forward Network Limitations

While the Universal Approximation Theorem suggests FFNs can approximate any function, gradient descent in LLMs typically leads to Pattern Storage rather than Algorithmic Discovery.

- *Piecewise Linearity:* FFNs using ReLU activations implement piecewise linear functions. These are architecturally incapable of implementing exact non-linear operations like multiplication ($x \times y$) over unbounded domains.
- *Hierarchical Pattern Assembly:* Rather than executing a generalizable algorithm, the model coordinates multiple FFN layers to approximate results based on memorized templates from training data. Performance “cliffs” occur

when tasks transition from pattern recognition to genuine rule discovery (e.g., the ARC-AGI-2 benchmark).

3. The Instruction-Execution Disconnect

Next-token prediction treats algorithmic descriptions and execution traces as equivalent sequences, which geometrically separates them in the model’s latent space.

- *Separated Pathways:* t-SNE projections reveal that instructional knowledge (how to do a task) and executorial know-how (doing the task) occupy distinct representational regions.
- *Reversal Curse:* Models can identify that “Tom Cruise’s mother is Mary Lee Pfeiffer” but often fail to answer “Who is Mary Lee Pfeiffer’s son?” because they lack the “lifted reasoning” required to apply symmetric logical rules across arbitrary entities.

II. Case Studies in Logical and Relational Failures

Problem Type	Description	Observed Failure
Decimal Comparison	Comparing values like 9.9 and 9.11	Models claim 9.11 is larger because 11 > 9, ignoring decimal place value
The Alice Problem	“Alice has 4 sisters and 1 brother. How many sisters does her brother have?”	Performance collapses (near 0%) as models fail to rebind Alice to a “sister” role from the brother’s perspective
Tabular Reasoning	Computing averages or maximums from CSV data	Significant deficits in maintaining comparison states or performing consistent increments
Relational Inference	Multi-step family tree reasoning (e.g., “IsUncle”)	LLMs perform dramatically worse than specialized models (DLMs) 100,000x smaller

III. Compensatory Strategies and the Metacognitive Bottleneck

1. **Self-Scaffolding (Chain-of-Thought):** Models generate intermediate steps to convert parallel constraints into serial computation. However, as problem complexity increases (e.g., 10-digit multiplication), errors in individual steps or final summations lead to total failure.
2. **Tool Delegation:** Models recognize problem types and call external systems (e.g., a calculator or SQL engine). This is highly effective but requires the model to correctly identify when it is “out of its depth.”
3. **Hybrid Architectures:** Integrating dedicated symbolic modules (e.g., OccamLLM or Logic-LM) into the transformer framework to handle arithmetic or logic internally.

The Metacognitive Dependency: All workarounds require “metacognition” — the ability of the model to assess its own internal limitations. Current architectures lack reliable self-awareness, often exhibiting overconfidence on tasks they are architecturally incapable of solving.

IV. Multi-Modal Sensor Fusion in High-Speed Systems

1. Late Fusion Architecture

The modular multi-modal late fusion approach integrates heterogeneous detection inputs (LiDAR and RADAR) to track objects at speeds up to 270 km/h.

- *LiDAR Strengths:* High accuracy in measuring an object’s longitudinal and lateral position.
- *RADAR Strengths:* Ability to directly measure an object’s velocity, compensating for LiDAR’s positional noise.
- *Object Matching:* Utilizes the Hungarian method with Euclidean distance as a cost function to associate new detections with existing tracks.

2. Perception Delay Compensation

A critical challenge in high-speed racing is perception software latency (up to 325ms). At 233 km/h, an object moves 13.7 meters during the perception delay.

- *Backward-Forward Integration:* When a delayed sensor input is received, the system performs a backward search in the object storage to find the state at the sensor’s timestamp.
- *Update and Integration:* The EKF update is applied to that historic state, which is then forward integrated using a kinematic model (CTRV — Constant Turn Rate and Velocity) to provide a corrected estimate for the current timestamp.

3. Performance and Reliability Metrics

The EKF-based fusion method achieved position tracking residuals below 0.1m in real-world validation at the Indy Autonomous Challenge.

Metric	LiDAR Clustering (C)	RADAR Detection (R)	Fused System (EKF)
Precision	0.59	0.32	0.50
Position Accuracy	High (0.3m STD)	Low (3.0m STD)	Optimized/Balanced
Velocity Accuracy	N/A	High (0.2m/s STD)	Robust

V. Key Insights and Conclusions

- **Comprehension vs. Competence:** LLMs represent an inversion of biological systems; they achieve sophisticated explanatory “comprehension” without the underlying “competence” for execution.
- **Architectural Inevitability:** Computational Split-Brain Syndrome is not a result of insufficient data but is inherent to the transformer architecture’s reliance on next-token prediction and piecewise linear FFNs.
- **Interpretability Concerns:** Findings in mechanistic interpretability (e.g., “arithmetic circuits”) may actually be observing “bags of heuristics” — path-dependent statistical patterns rather than universal computational principles.
- **Systemic Reliability:** In safety-critical applications like autonomous racing, reliability is derived from explicit architectural support for symbolic operations (e.g., CTRV models and EKF) rather than emergent pattern completion.

Future AI development must address these gaps through metacognitive scaffolding, lifted representations, and architecturally grounded execution to move beyond sophisticated pattern matching toward genuine generalizable intelligence.

A.2 Step 2 Output — Claude Draft (Full Text)

Comprehension Without Competence: Architectural Limits in LLM Reasoning and Lessons from Autonomous Sensor Fusion

Executive Summary

Contemporary AI research reveals a striking dichotomy between linguistic fluency and execution reliability. Large Language Models (LLMs) exhibit what can be termed “Computational Split-Brain Syndrome”: the capacity to articulate complex procedural rules while failing to execute them, rooted in architectural limitations in symbol representation and feed-forward network (FFN) computation. This failure mode stems from a core property of next-token prediction objectives, which decouple instructional knowledge from execution pathways.

By contrast, the field of autonomous robotics achieves reliability through explicit architectural modularity. Techniques such as Extended Kalman Filters (EKF) and multi-modal late fusion — combining LiDAR and RADAR — enable high-speed object tracking, while “delay compensation” using backward-forward integration reconciles perception latency with real-time motion planning. Together, these two domains illustrate a broader principle: systems that rely on emergent pattern completion struggle with symbolic precision, while systems built on explicit mathematical frameworks achieve it reliably.

I. Architectural Limits of LLMs in Symbolic Reasoning

Computational Split-Brain Syndrome describes a systematic dissociation in which LLMs can perfectly explain a principle while being unable to reliably execute it. This dissociation traces to three distinct architectural constraints.

1. *Representation: The Failure of Domain Binding* — LLM token embeddings encode context-weighted averages that resist “automatic domain binding.” Because the training objective optimizes a single embedding across many contexts, tokens accumulate conflicting semantic associations rather than clean, domain-specific meanings. A token such as “9.11” carries associations with historical dates and software versioning alongside its numerical value, so the model fails to strip these away when comparing 9.11 to 9.9. Embedding analysis also shows that numerical tokens do not preserve isometric properties such as numerical order or distance — in certain latent spaces, “10” may sit closer to “1” than to “9.”

2. *Computation: Feed-Forward Network Limitations* — While the Universal Approximation Theorem suggests FFNs can in principle approximate any function, gradient descent in practice tends to produce pattern storage rather than algorithmic discovery. FFNs using ReLU activations implement piecewise linear functions, which are architecturally incapable of implementing exact non-linear operations such as multiplication over unbounded domains. Rather than executing a generalizable algorithm, models coordinate multiple FFN layers to approximate results using memorized templates, producing performance “cliffs” on benchmarks like ARC-AGI-2.

3. *The Instruction-Execution Disconnect* — Because next-token prediction treats algorithmic descriptions and execution traces as equivalent sequences, the two are geometrically separated in the model’s latent space. t-SNE projections show that instructional knowledge and executorial know-how occupy distinct representational regions. This produces the “Reversal Curse,” in which a model may state that “Tom Cruise’s mother is Mary Lee Pfeiffer” yet fail to answer “Who is Mary Lee Pfeiffer’s son?” — lacking the “lifted reasoning” needed to apply symmetric logical rules across arbitrary entities.

These three constraints — representation, computation, and instruction-execution separation — compound to produce the case-level failures documented next.

II. Case Studies in Logical and Relational Failures

Problem Type	Description	Observed Failure
Decimal Comparison	Comparing values like 9.9 and 9.11	Models claim 9.11 is larger because 11 > 9, ignoring decimal place value
The Alice Problem	“Alice has 4 sisters and 1 brother. How many sisters does her brother have?”	Performance collapses to near 0% as models fail to rebind Alice to a “sister” role from the brother’s perspective

Tabular Reasoning	Computing averages or maximums from CSV data	Significant deficits in maintaining comparison states or performing consistent increments
Relational Inference	Multi-step family tree reasoning (e.g., “IsUncle”)	LLMs perform dramatically worse than specialized models (DLMs) 100,000x smaller

Each case reflects the same underlying issue: the model can describe the relevant rule but cannot reliably bind and manipulate the symbols the rule depends on.

III. Compensatory Strategies and the Metacognitive Bottleneck

- **Self-Scaffolding (Chain-of-Thought):** Models generate intermediate steps to convert parallel constraints into serial computation. As problem complexity increases — for example, 10-digit multiplication — errors in individual steps or in final summation compound into total failure.
- **Tool Delegation:** Models recognize a problem type and call an external system, such as a calculator or SQL engine. This is highly effective, but depends on the model correctly identifying when it is “out of its depth.”
- **Hybrid Architectures:** Dedicated symbolic modules — such as OccamLLM or Logic-LM — are integrated into the transformer framework to handle arithmetic or logic internally.

All three strategies depend on metacognition — the model’s ability to assess its own internal limitations. Current architectures lack reliable self-awareness of this kind, and frequently exhibit overconfidence on tasks they are architecturally incapable of solving. This dependency is the central obstacle to closing the gap between comprehension and competence.

IV. Multi-Modal Sensor Fusion in High-Speed Systems

- 1. Late Fusion Architecture* — A modular multi-modal late fusion approach integrates heterogeneous detection inputs from LiDAR and RADAR to track objects at speeds up to 270 km/h. LiDAR offers high accuracy in measuring an object’s longitudinal and lateral position, while RADAR directly measures velocity, compensating for LiDAR’s positional noise. The Hungarian method, using Euclidean distance as a cost function, associates new detections with existing tracks.
- 2. Perception Delay Compensation* — A critical challenge in high-speed racing is perception software latency, which can reach up to 325ms; at 233 km/h, an object moves 13.7 meters during this delay window. When a delayed sensor input arrives, the system performs a backward search through object storage to find the state at the sensor’s original timestamp. The EKF update is applied to that historic state, then forward-integrated using a Constant Turn Rate and Velocity (CTRV) kinematic model to produce a corrected estimate for the current timestamp.
- 3. Performance and Reliability Metrics* — The EKF-based fusion method achieved position tracking residuals below 0.1m in real-world validation at the Indy Autonomous Challenge.

Metric	LiDAR Clustering (C)	RADAR Detection (R)	Fused System (EKF)
Precision	0.59	0.32	0.50
Position Accuracy	High (0.3m STD)	Low (3.0m STD)	Optimized/Balanced
Velocity Accuracy	N/A	High (0.2m/s STD)	Robust

The fused system does not simply average its inputs — it selectively leverages each sensor’s strength (LiDAR’s positional precision, RADAR’s velocity precision) while compensating for the other’s weaknesses, yielding a more balanced and robust result than either sensor alone.

V. Key Insights and Conclusions

- **Comprehension vs. Competence:** LLMs invert the typical relationship seen in biological systems — they achieve sophisticated explanatory “comprehension” without the underlying “competence” required for execution.
- **Architectural Inevitability:** Computational Split-Brain Syndrome is not a byproduct of insufficient training data; it is inherent to the transformer architecture’s reliance on next-token prediction and piecewise linear FFNs.
- **Interpretability Concerns:** Findings in mechanistic interpretability — such as identified “arithmetic circuits” — may in fact be observing “bags of heuristics”: path-dependent statistical patterns rather than universal computational principles.
- **Systemic Reliability:** In safety-critical applications like autonomous racing, reliability is derived from explicit architectural support for symbolic operations — CTRV models and EKF — rather than from emergent pattern completion.

Taken together, these findings point to a clear path forward for AI development: closing the gap between comprehension and competence will require metacognitive scaffolding, lifted representations, and architecturally grounded execution — moving beyond sophisticated pattern matching toward genuine, generalizable intelligence.

A.3 Step 3 Output — ChatGPT Critique Log (Full Text)

Based on the two outputs, only Output 1 was received in the initial message; Output 2 was initially reported missing, with the attached file containing the same content as Output 1.

Critique:

- *Accuracy:* Could not compare the draft against the outline while Output 2 (the draft) was missing. Based on Output

- 1 alone, the outline was assessed as comprehensive, logically structured, and covering all major concepts.
- *Tone & Clarity*: The writing was already professional and technically precise. Suggested improvements included shortening a few long sentences, reducing repetition (e.g., repeated mentions of “architectural limitations”), and smoothing transitions between sections.

A final revised version could not be produced until Output 2 (the draft) was supplied. Once supplied, the promised deliverable was:

1. A 2-bullet critique (Accuracy; Tone & Technical Clarity).
2. A final revised version with all missing constraints restored, direct and scannable wording, passive voice removed where appropriate, standardized technical terminology, and a publication-ready final draft preserving the original meaning.